


SVHN Digit Recognition

Neural Networks, Tensorflow, Keras

'Femi Bolarinwa

Digit Image Snapshot



- Images available in grayscale stored in stack of 32 x 32 pixels
- 42,000 and 18,000 images used for training and testing respectively
- Pixels normalized to prevent exploding gradient
- One-hot encoding for target variables (image classifications)

Model 1

Artificial Neural Network

- First layer: 64 neurons, RELU activation
- Second layer: 32 neurons, RELU activation
- Output layer: 10 neurons, softmax activation
- 68,010 trainable parameters (weights & biases)

```
model_1.summary()
```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 64)	65600
dense_4 (Dense)	(None, 32)	2080
dense_5 (Dense)	(None, 10)	330

```
=====  
Total params: 68,010
```


```
Trainable params: 68,010
```

```
Non-trainable params: 0
```

Model 1 (ANN)

Model Training

- 20 epochs
- 'Adam' optimizer, learning rate = 0.001
- Categorical-crossentropy loss function, accuracy metric
- ~70% accuracy on training and validation
- Decent accuracy



Model 2

Artificial Neural Network

- 6 layers, 554 neurons
- Dropout & BatchNormalization layers for regularization
- 310,186 trainable parameters (weights and biases)

```
model_2.summary()
```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
dense_6 (Dense)	(None, 256)	262400
dense_7 (Dense)	(None, 128)	32896
dropout_1 (Dropout)	(None, 128)	0
dense_8 (Dense)	(None, 64)	8256
dense_9 (Dense)	(None, 64)	4160
dense_10 (Dense)	(None, 32)	2080
batch_normalization_1 (Batch Normalization)	(None, 32)	128
dense_11 (Dense)	(None, 10)	330

```
Total params: 310,250
```


```
Trainable params: 310,186
```

```
Non-trainable params: 64
```

Model 2 (ANN)

Model Training

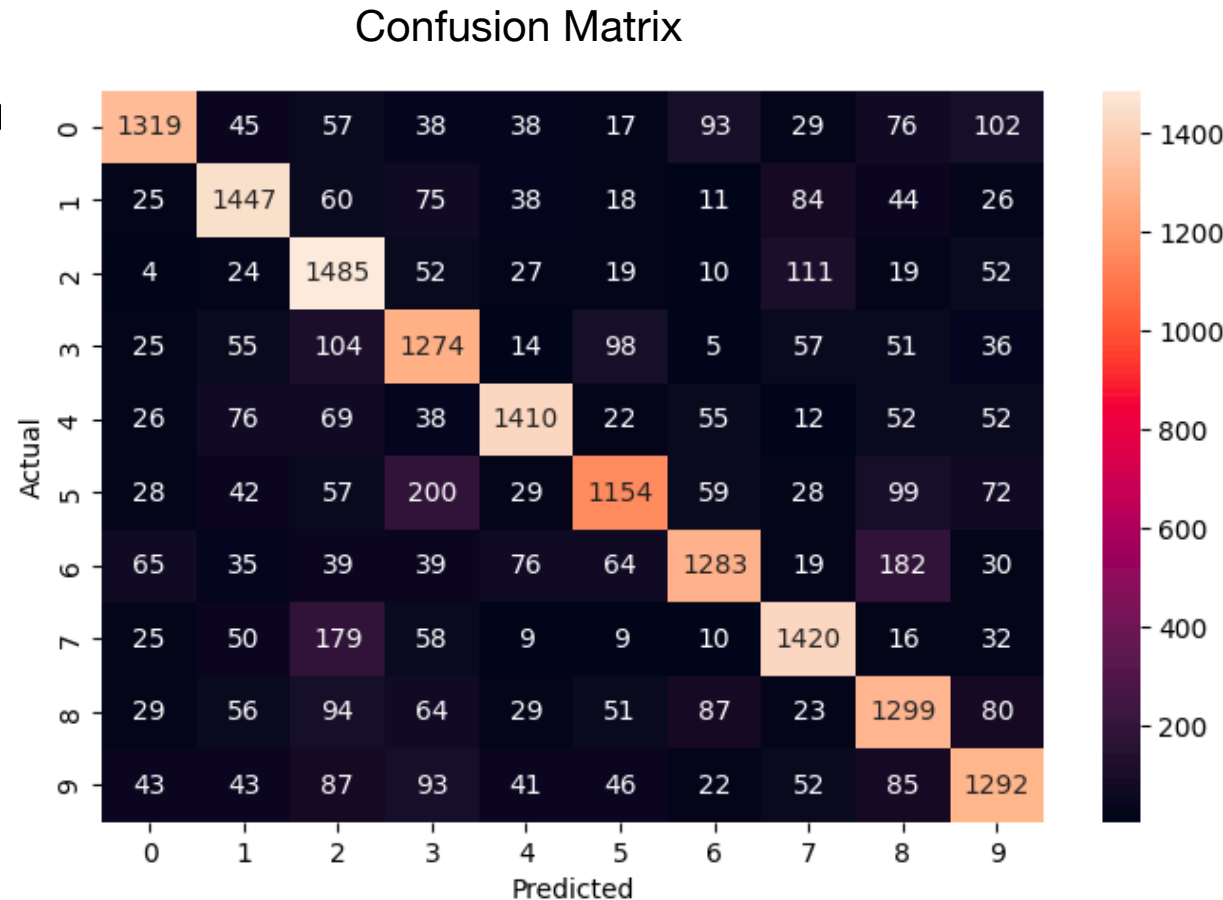
- 30 epochs
- 'Adam' optimizer, learning rate = 0.0005
- Categorical-crossentropy loss function, accuracy metric
- ~74% accuracy on training and validation
- Bigger architecture, but no significant accuracy boost



Model 2 (ANN)

Evaluation on Unseen Test data

- Accuracy: 74%
- Recall: 74%
- Precision: 75%
- f1-score: 74%
- Similar accuracy on training and test dataset.
- Decent model, not overfitted



Model 3

Convolutional Neural Network

- 6 layers
- 4 convolution filters for feature extraction
- Padding and MaxPooling
- Dropout & BatchNormalization layers for regularization
- 164,170 trainable parameters (weights and biases)

```
cnn_model_2.summary()

Model: "sequential_2"

```

Layer (type)	Output Shape	Param #
conv2d_8 (Conv2D)	(None, 32, 32, 16)	160
leaky_re_lu_10 (LeakyReLU)	(None, 32, 32, 16)	0
conv2d_9 (Conv2D)	(None, 32, 32, 32)	4640
leaky_re_lu_11 (LeakyReLU)	(None, 32, 32, 32)	0
max_pooling2d_4 (MaxPooling 2D)	(None, 16, 16, 32)	0
batch_normalization_4 (Batch Normalization)	(None, 16, 16, 32)	128
conv2d_10 (Conv2D)	(None, 16, 16, 32)	9248
leaky_re_lu_12 (LeakyReLU)	(None, 16, 16, 32)	0
conv2d_11 (Conv2D)	(None, 16, 16, 64)	18496
leaky_re_lu_13 (LeakyReLU)	(None, 16, 16, 64)	0
max_pooling2d_5 (MaxPooling 2D)	(None, 8, 8, 64)	0
batch_normalization_5 (Batch Normalization)	(None, 8, 8, 64)	256
flatten_2 (Flatten)	(None, 4096)	0
dense_4 (Dense)	(None, 32)	131104
leaky_re_lu_14 (LeakyReLU)	(None, 32)	0
dropout_2 (Dropout)	(None, 32)	0
dense_5 (Dense)	(None, 10)	330

```


Total params: 164,362
Trainable params: 164,170
Non-trainable params: 192

```


CNN Model

Model Training


- 20 epochs
- 'Adam' optimizer, learning_rate = 0.001
- Categorical-crossentropy loss function, accuracy metric
- >90% accuracy on training and validation



CNN Model

Evaluation on Unseen Test data

- Accuracy: 91%
- Recall: 91%
- Precision: 91%
- f1-score: 91%
- Similar accuracy on training and test dataset.
- Significantly less misclassification
- Great model, not overfitted



Insight and Recommendation

- CNN appears to be better than ANN for feature extraction and hence higher accuracy
- However, this comes at the cost of more computational capacity requirement
- There is always a need to balance the need for higher accuracy and cost of complex architecture and hence computation time